

titanic-prediction-model

June 27, 2024

Titanic Survival Prediction Model

By Elma Fortunate Phiri

About Dataset:

The sinking of the Titanic is one of the most infamous shipwrecks in history.

On April 15, 1912, during her maiden voyage, the widely considered “unsinkable” RMS Titanic sank after colliding with an iceberg. Unfortunately, there weren’t enough lifeboats for everyone onboard, resulting in the death of 1502 out of 2224 passengers and crew.

While there was some element of luck involved in surviving, it seems some groups of people were more likely to survive than others.

Abstract: To perform data cleaning and exploratory data analysis (EDA) on the dataset. Explore the relationships between variables and identify patterns and trends in the data

Aim: To build a predictive model that answers the question: “what sorts of people were more likely to survive?” using passenger data (ie name, age, gender, socio-economic class, etc).

Data collection and Preprocessing

Importing Python Libraries

```
[269]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import ydata_profiling as pp
%matplotlib inline

[344]: from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split, RandomizedSearchCV,
↳GridSearchCV
from sklearn.metrics import accuracy_score, confusion_matrix,
↳classification_report, precision_score, RocCurveDisplay
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from scipy.stats import randint
```

```
[271]: titanic = pd.read_csv("C:\\Users\\elmaf\\Desktop\\ProgidyInfoTech\\Titanic_
dataset\\titanic3.csv")
titanic
```

```
[271]:
```

	pclass	survived	name \
0	1.0	1.0	Allen, Miss. Elisabeth Walton
1	1.0	1.0	Allison, Master. Hudson Trevor
2	1.0	0.0	Allison, Miss. Helen Loraine
3	1.0	0.0	Allison, Mr. Hudson Joshua Creighton
4	1.0	0.0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)
...
1305	3.0	0.0	Zabour, Miss. Thamine
1306	3.0	0.0	Zakarian, Mr. Mapriededer
1307	3.0	0.0	Zakarian, Mr. Ortin
1308	3.0	0.0	Zimmerman, Mr. Leo
1309	NaN	NaN	NaN

	sex	age	sibsp	parch	ticket	fare	cabin	embarked	boat \
0	female	29.0000	0.0	0.0	24160	211.3375	B5	S	2
1	male	0.9167	1.0	2.0	113781	151.5500	C22 C26	S	11
2	female	2.0000	1.0	2.0	113781	151.5500	C22 C26	S	NaN
3	male	30.0000	1.0	2.0	113781	151.5500	C22 C26	S	NaN
4	female	25.0000	1.0	2.0	113781	151.5500	C22 C26	S	NaN
...
1305	female	NaN	1.0	0.0	2665	14.4542	NaN	C	NaN
1306	male	26.5000	0.0	0.0	2656	7.2250	NaN	C	NaN
1307	male	27.0000	0.0	0.0	2670	7.2250	NaN	C	NaN
1308	male	29.0000	0.0	0.0	315082	7.8750	NaN	S	NaN
1309	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

	body	home.dest
0	NaN	St Louis, MO
1	NaN	Montreal, PQ / Chesterville, ON
2	NaN	Montreal, PQ / Chesterville, ON
3	135.0	Montreal, PQ / Chesterville, ON
4	NaN	Montreal, PQ / Chesterville, ON
...
1305	NaN	NaN
1306	304.0	NaN
1307	NaN	NaN
1308	NaN	NaN
1309	NaN	NaN

[1310 rows x 14 columns]

Data Exploration

```
[272]: titanic.head()
```

```
[272]:
```

	pclass	survived		name	sex	\
0	1.0	1.0		Allen, Miss. Elisabeth Walton	female	
1	1.0	1.0		Allison, Master. Hudson Trevor	male	
2	1.0	0.0		Allison, Miss. Helen Loraine	female	
3	1.0	0.0		Allison, Mr. Hudson Joshua Creighton	male	
4	1.0	0.0		Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	female	

	age	sibsp	parch	ticket	fare	cabin	embarked	boat	body	\
0	29.0000	0.0	0.0	24160	211.3375	B5	S	2	NaN	
1	0.9167	1.0	2.0	113781	151.5500	C22 C26	S	11	NaN	
2	2.0000	1.0	2.0	113781	151.5500	C22 C26	S	NaN	NaN	
3	30.0000	1.0	2.0	113781	151.5500	C22 C26	S	NaN	135.0	
4	25.0000	1.0	2.0	113781	151.5500	C22 C26	S	NaN	NaN	

	home.dest
0	St Louis, MO
1	Montreal, PQ / Chesterville, ON
2	Montreal, PQ / Chesterville, ON
3	Montreal, PQ / Chesterville, ON
4	Montreal, PQ / Chesterville, ON

```
[273]: titanic.shape
```

```
[273]: (1310, 14)
```

```
[274]: titanic.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1310 entries, 0 to 1309
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   pclass      1309 non-null   float64
1   survived    1309 non-null   float64
2   name        1309 non-null   object
3   sex         1309 non-null   object
4   age         1046 non-null   float64
5   sibsp       1309 non-null   float64
6   parch       1309 non-null   float64
7   ticket      1309 non-null   object
8   fare        1308 non-null   float64
9   cabin       295 non-null    object
10  embarked    1307 non-null   object
11  boat        486 non-null    object
12  body        121 non-null    float64
```

```
13 home.dest 745 non-null object
dtypes: float64(7), object(7)
memory usage: 143.4+ KB
```

```
[275]: titanic.describe()
```

```
[275]:
```

	pclass	survived	age	sibsp	parch \
count	1309.000000	1309.000000	1046.000000	1309.000000	1309.000000
mean	2.294882	0.381971	29.881135	0.498854	0.385027
std	0.837836	0.486055	14.413500	1.041658	0.865560
min	1.000000	0.000000	0.166700	0.000000	0.000000
25%	2.000000	0.000000	21.000000	0.000000	0.000000
50%	3.000000	0.000000	28.000000	0.000000	0.000000
75%	3.000000	1.000000	39.000000	1.000000	0.000000
max	3.000000	1.000000	80.000000	8.000000	9.000000

	fare	body
count	1308.000000	121.000000
mean	33.295479	160.809917
std	51.758668	97.696922
min	0.000000	1.000000
25%	7.895800	72.000000
50%	14.454200	155.000000
75%	31.275000	256.000000
max	512.329200	328.000000

```
[276]: titanic.columns
```

```
[276]: Index(['pclass', 'survived', 'name', 'sex', 'age', 'sibsp', 'parch', 'ticket',  
          'fare', 'cabin', 'embarked', 'boat', 'body', 'home.dest'],  
          dtype='object')
```

```
[277]: titanic.index
```

```
[277]: RangeIndex(start=0, stop=1310, step=1)
```

```
[278]: pp.ProfileReport(titanic)
```

```
Summarize dataset: 0%|          | 0/5 [00:00<?, ?it/s]
Generate report structure: 0%|          | 0/1 [00:00<?, ?it/s]
Render HTML: 0%|          | 0/1 [00:00<?, ?it/s]
<IPython.core.display.HTML object>
```

```
[278]:
```

The Profile report highlighted the following: 1. We have 14 variables in our dataset 2. We have 1310 observations 3. Age has 20.2% missing values 5. Cabin has 687 (77.5%) missing values 6.

Boat has 62.9% missing values 7. Body has 90.8% missing values 9. home.dest has 565 (43.1%) missing values

```
[279]: titanic.dtypes
```

```
[279]: pclass      float64
survived      float64
name          object
sex           object
age           float64
sibsp         float64
parch         float64
ticket        object
fare          float64
cabin         object
embarked      object
boat          object
body          float64
home.dest     object
dtype: object
```

Age has wrong data dtype. It should be int but its float

```
[280]: titanic["sex"].value_counts()
```

```
[280]: male      843
female    466
Name: sex, dtype: int64
```


Data Preprocessing

```
[281]: #Checking for null values
titanic.isnull().sum().sort_values(ascending=False)
```

```
[281]: body      1189
cabin     1015
boat       824
home.dest  565
age        264
embarked     3
fare         2
pclass       1
survived     1
name         1
sex          1
sibsp        1
parch        1
ticket       1
```

dtype: int64

```
[282]: #Visualizing missing values
sns.heatmap(titanic.isnull())
plt.show()
```



```
[283]: #dropping columns with more than 40% missing values
titanic = titanic.drop(["cabin", 'body','boat','home.dest'], axis=1)
titanic.head()
```

```
[283]:
```

	pclass	survived	name	sex	\
0	1.0	1.0	Allen, Miss. Elisabeth Walton	female	
1	1.0	1.0	Allison, Master. Hudson Trevor	male	
2	1.0	0.0	Allison, Miss. Helen Loraine	female	
3	1.0	0.0	Allison, Mr. Hudson Joshua Creighton	male	
4	1.0	0.0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	female	

	age	sibsp	parch	ticket	fare	embarked
--	-----	-------	-------	--------	------	----------

0	29.0000	0.0	0.0	24160	211.3375	S
1	0.9167	1.0	2.0	113781	151.5500	S
2	2.0000	1.0	2.0	113781	151.5500	S
3	30.0000	1.0	2.0	113781	151.5500	S
4	25.0000	1.0	2.0	113781	151.5500	S

```
[284]: #Renaming columns
titanic.rename(columns={
    " pclass": "Pclass",
    "survived": "Survived",
    "name": "Name",
    "sex": "Sex",
    "age": "Age",
    "sibsp": "No of Sibilings/ Spouse aboard",
    "parch": "No of Parents/Children",
    "ticket": "Ticket Number",
    "fare": "Fare",
    "embarked": "Embarked"
}, inplace=True)
titanic.head()
```

```
[284]:
```

	pclass	Survived	Name	Sex \
0	1.0	1.0	Allen, Miss. Elisabeth Walton	female
1	1.0	1.0	Allison, Master. Hudson Trevor	male
2	1.0	0.0	Allison, Miss. Helen Loraine	female
3	1.0	0.0	Allison, Mr. Hudson Joshua Creighton	male
4	1.0	0.0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	female

	Age	No of Sibilings/ Spouse aboard	No of Parents/Children \
0	29.0000	0.0	0.0
1	0.9167	1.0	2.0
2	2.0000	1.0	2.0
3	30.0000	1.0	2.0
4	25.0000	1.0	2.0

	Ticket Number	Fare	Embarked
0	24160	211.3375	S
1	113781	151.5500	S
2	113781	151.5500	S
3	113781	151.5500	S
4	113781	151.5500	S

```
[285]: titanic["Embarked"].value_counts()
```

```
[285]: S    914
      C    270
      Q    123
```

Name: Embarked, dtype: int64

```
[286]: #Replace values in embarked column
titanic["Embarked"] = titanic["Embarked"].replace({"C": "Cherbourg", "Q": "Queenstown", "S": "Southampton"})
titanic.sample(10)
```

```
[286]:
```

	pclass	Survived	Name	Sex	\
1147	3.0	0.0	Riihivouri, Miss. Susanna Juhantytar "Sanni"	female	
719	3.0	1.0	Cohen, Mr. Gurshon "Gus"	male	
469	2.0	1.0	Keane, Miss. Nora A	female	
195	1.0	1.0	Maioni, Miss. Roberta	female	
384	2.0	0.0	Cunningham, Mr. Alfred Fleming	male	
1157	3.0	0.0	Rosblom, Mr. Viktor Richard	male	
847	3.0	0.0	Hanna, Mr. Mansour	male	
771	3.0	1.0	Devaney, Miss. Margaret Delia	female	
289	1.0	1.0	Taussig, Miss. Ruth	female	
414	2.0	0.0	Gale, Mr. Shadrach	male	

	Age	No of Sibilings/ Spouse aboard	No of Parents/Children	\
1147	22.0	0.0	0.0	
719	18.0	0.0	0.0	
469	NaN	0.0	0.0	
195	16.0	0.0	0.0	
384	NaN	0.0	0.0	
1157	18.0	1.0	1.0	
847	23.5	0.0	0.0	
771	19.0	0.0	0.0	
289	18.0	0.0	2.0	
414	34.0	1.0	0.0	


	Ticket Number	Fare	Embarked
1147	3101295	39.6875	Southampton
719	A/5 3540	8.0500	Southampton
469	226593	12.3500	Queenstown
195	110152	86.5000	Southampton
384	239853	0.0000	Southampton
1157	370129	20.2125	Southampton
847	2693	7.2292	Cherbourg
771	330958	7.8792	Queenstown
289	110413	79.6500	Southampton
414	28664	21.0000	Southampton

```
[287]: #Checking for null values
titanic.isnull().sum().sort_values(ascending=False)
```



```
[287]: Age                264
      Embarked           3
      Fare               2
      pclass             1
      Survived           1
      Name               1
      Sex                1
      No of Sibilings/ Spouse aboard 1
      No of Parents/Children 1
      Ticket Number      1
      dtype: int64
```

```
[288]: #Histogram for Age column
      sns.histplot(x = "Age", data = titanic)
      plt.title("Age distribution")
      plt.ylabel("Number of People")
      plt.show()
```



```
[289]: #Handling Missing data in Age column by filling na values
      Age_mean = titanic["Age"].mean()
```

```
Age_mean
titanic['Age']=titanic['Age'].fillna(Age_mean)
```

```
[290]: #Handling Missing values in Embarked column by filling na values
titanic["Embarked"] = titanic["Embarked"].fillna(titanic["Embarked"].
↪value_counts().idxmax())
```

```
[291]: #Handling Missing values by filling na values
titanic["Fare"] = titanic["Fare"].fillna(titanic["Fare"].mean())
```

```
[292]: #Checking for null values
titanic.isnull().sum().sort_values(ascending=False)
```

```
[292]: pclass                1
Survived                  1
Name                      1
Sex                      1
No of Sibilings/ Spouse aboard  1
No of Parents/Children      1
Ticket Number             1
Age                       0
Fare                      0
Embarked                  0
dtype: int64
```

```
[293]: #drop na values
titanic.dropna(inplace= True)
```

```
[294]: #Checking for null values
titanic.isnull().sum().sort_values(ascending=False)
```

```
[294]: pclass                0
Survived                  0
Name                      0
Sex                      0
Age                      0
No of Sibilings/ Spouse aboard  0
No of Parents/Children      0
Ticket Number             0
Fare                      0
Embarked                  0
dtype: int64
```

```
[295]: titanic.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1309 entries, 0 to 1308
```

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	pclass	1309 non-null	float64
1	Survived	1309 non-null	float64
2	Name	1309 non-null	object
3	Sex	1309 non-null	object
4	Age	1309 non-null	float64
5	No of Sibilings/ Spouse aboard	1309 non-null	float64
6	No of Parents/Children	1309 non-null	float64
7	Ticket Number	1309 non-null	object
8	Fare	1309 non-null	float64
9	Embarked	1309 non-null	object

dtypes: float64(6), object(4)

memory usage: 112.5+ KB

```
[296]: #changing datatypes
titanic=titanic.astype({
    "Age": "int64",
    "pclass": "int64",
    "Survived": "int64"
})
titanic.info()
```

<class 'pandas.core.frame.DataFrame'>

Int64Index: 1309 entries, 0 to 1308

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	pclass	1309 non-null	int64
1	Survived	1309 non-null	int64
2	Name	1309 non-null	object
3	Sex	1309 non-null	object
4	Age	1309 non-null	int64
5	No of Sibilings/ Spouse aboard	1309 non-null	float64
6	No of Parents/Children	1309 non-null	float64
7	Ticket Number	1309 non-null	object
8	Fare	1309 non-null	float64
9	Embarked	1309 non-null	object

dtypes: float64(3), int64(3), object(4)

memory usage: 112.5+ KB

```
[297]: titanic["No of Parents/Children"] = titanic["No of Parents/Children"].
        ↪astype("int64")
titanic['No of Sibilings/ Spouse aboard'] = titanic['No of Sibilings/ Spouse_
        ↪aboard'].astype("int64")
```

```
titanic.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1309 entries, 0 to 1308
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   pclass                                1309 non-null   int64
1   Survived                             1309 non-null   int64
2   Name                                  1309 non-null   object
3   Sex                                   1309 non-null   object
4   Age                                  1309 non-null   int64
5   No of Sibilings/ Spouse aboard        1309 non-null   int64
6   No of Parents/Children                1309 non-null   int64
7   Ticket Number                         1309 non-null   object
8   Fare                                  1309 non-null   float64
9   Embarked                             1309 non-null   object
dtypes: float64(1), int64(5), object(4)
memory usage: 112.5+ KB
```

```
[298]: titanic.head()
```

```
[298]:
```

	pclass	Survived	Name	Sex	\
0	1	1	Allen, Miss. Elisabeth Walton	female	
1	1	1	Allison, Master. Hudson Trevor	male	
2	1	0	Allison, Miss. Helen Loraine	female	
3	1	0	Allison, Mr. Hudson Joshua Creighton	male	
4	1	0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	female	

	Age	No of Sibilings/ Spouse aboard	No of Parents/Children	Ticket	Number	\
0	29	0		0	24160	
1	0	1		2	113781	
2	2	1		2	113781	
3	30	1		2	113781	
4	25	1		2	113781	

	Fare	Embarked
0	211.3375	Southampton
1	151.5500	Southampton
2	151.5500	Southampton
3	151.5500	Southampton
4	151.5500	Southampton

```
[299]: #Correlation matrix
corr_matrix = titanic.corr()
corr_matrix
```

```
C:\Users\elmaf\AppData\Local\Temp\ipykernel_16020\2987745338.py:2:
```

FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
corr_matrix = titanic.corr()
```

[299]:

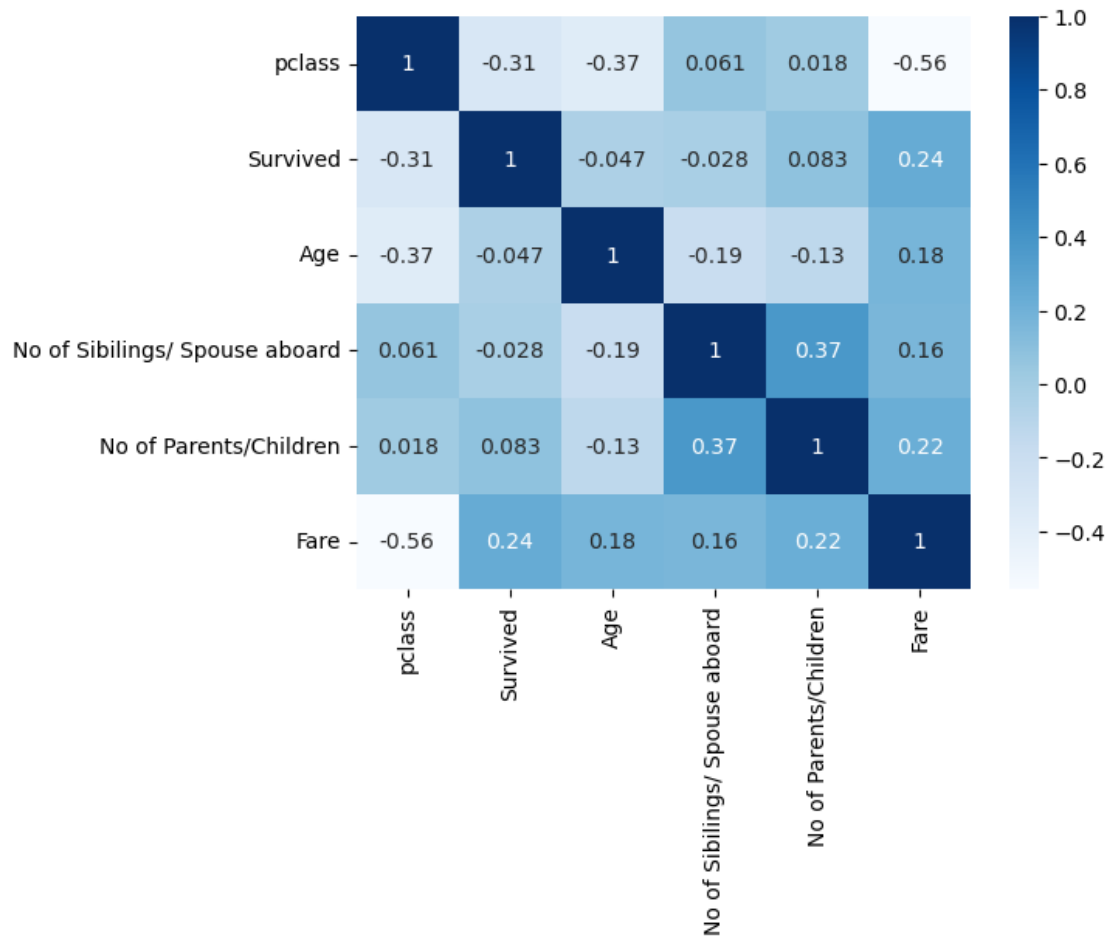
	pclass	Survived	Age	\
pclass	1.000000	-0.312469	-0.372115	
Survived	-0.312469	1.000000	-0.047021	
Age	-0.372115	-0.047021	1.000000	
No of Sibilings/ Spouse aboard	0.060832	-0.027825	-0.190345	
No of Parents/Children	0.018322	0.082660	-0.128821	
Fare	-0.558477	0.244208	0.175114	

	No of Sibilings/ Spouse aboard	\
pclass	0.060832	
Survived	-0.027825	
Age	-0.190345	
No of Sibilings/ Spouse aboard	1.000000	
No of Parents/Children	0.373587	
Fare	0.160224	

	No of Parents/Children	Fare
pclass	0.018322	-0.558477
Survived	0.082660	0.244208
Age	-0.128821	0.175114
No of Sibilings/ Spouse aboard	0.373587	0.160224
No of Parents/Children	1.000000	0.221522
Fare	0.221522	1.000000

[300]:

```
#Visualising our correlation matrix
sns.heatmap(corr_matrix, annot= True, cmap= 'Blues')
plt.show()
```

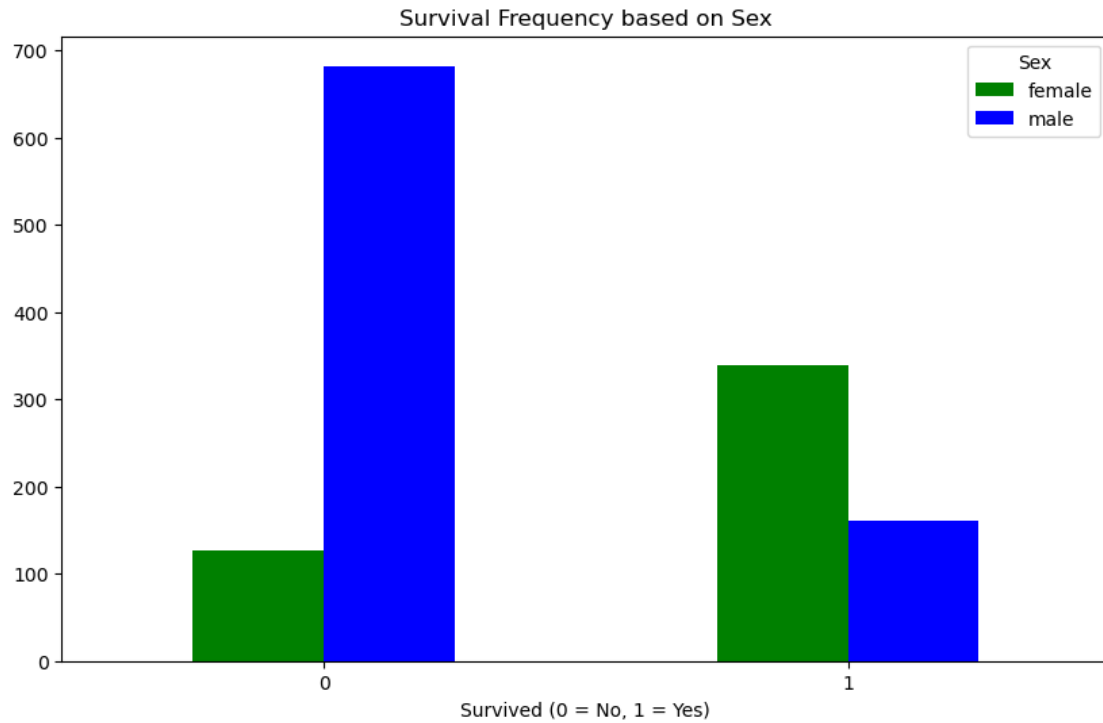


```
[301]: #Survival frequency according to gender
pd.crosstab(titanic.Survived , titanic.Sex)
```

```
[301]: Sex      female  male
Survived
0         127    682
1         339    161
```

```
[302]: pd.crosstab(titanic.Survived , titanic.Sex).plot(kind = 'bar', figsize=(10,6),
↳ color = ['green','blue'])
plt.xlabel("Survived (0 = No, 1 = Yes)")
plt.title("Survival Frequency based on Sex")
plt.xticks(rotation=0)
```

```
[302]: (array([0, 1]), [Text(0, 0, '0'), Text(1, 0, '1')])
```



```
[303]: #Survived frequency according to age
survival_age = pd.crosstab(titanic.Survived , titanic.Age)
survival_age
```

```
[303]: Age      0   1   2   3   4   5   6   7   8   9   ...  63  64  65  66  67  70  \
Survived
0          2   3   8   2   3   1   3   2   2   6   ...  2   3   3   1   1   3
1         10   7   4   5   7   4   3   2   4   4   ...  2   2   0   0   0   0

Age      71  74  76  80
Survived
0          2   1   0   0
1          0   0   1   1

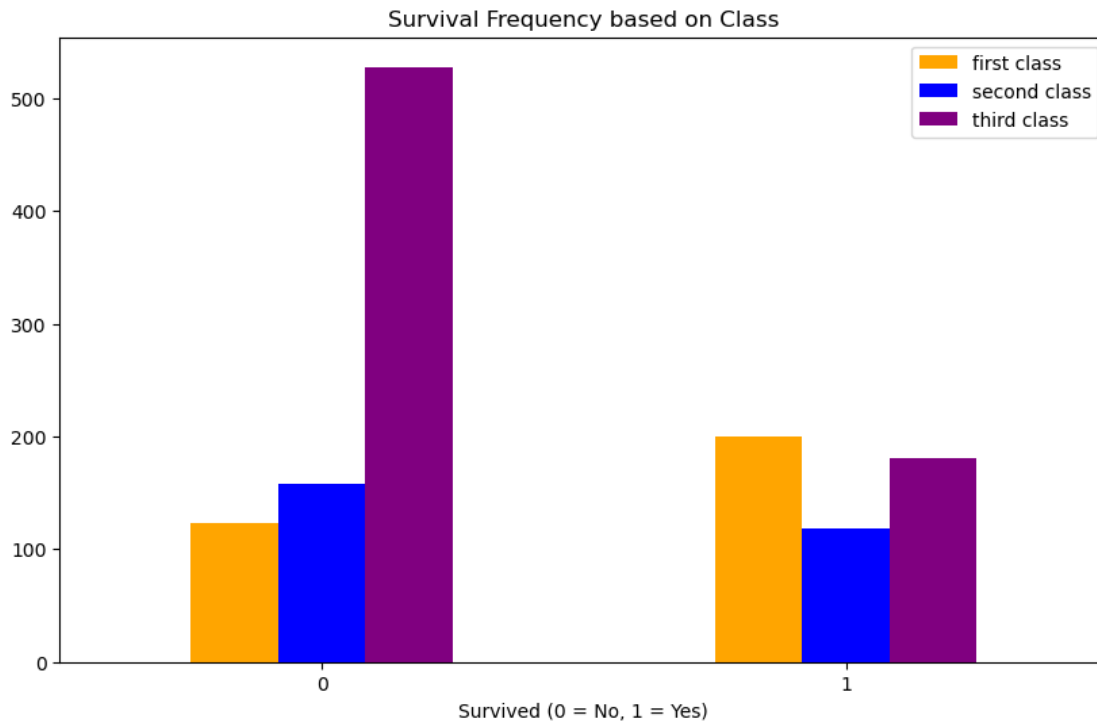
[2 rows x 73 columns]
```

```
[304]: #Survival frequency based on class
pd.crosstab(titanic.Survived, titanic.pclass)
```

```
[304]: pclass      1     2     3
Survived
0       123   158   528
1       200   119   181
```

```
[305]: #Visualizing survival frequency based on class
crosstab = pd.crosstab(titanic.Survived, titanic.pclass)
crosstab.plot(kind='bar', figsize=(10,6), color=['orange', 'blue', 'purple'])
plt.legend(["first class", "second class", "third class"])
plt.xlabel("Survived (0 = No, 1 = Yes)")
plt.title("Survival Frequency based on Class")
plt.xticks(rotation=0)
```

```
[305]: (array([0, 1]), [Text(0, 0, '0'), Text(1, 0, '1')])
```




```
[306]: #Survival frequency based on embarking station
pd.crosstab(titanic.Survived, titanic.Embarked)
```

```
[306]: Embarked  Cherbourg  Queenstown  Southampton
Survived
0           120         79           610
1           150         44           306
```


```
[307]: #Visualising the #Survival frequency based on embarking station
pd.crosstab(titanic.Survived, titanic.Embarked).plot(kind = "bar", color=□
↳ ["gold", "skyblue", "pink"])
plt.title("Survival Frequency base on Embarking station")
plt.xlabel("(Survived: 0 = No, 1 = Yes)")
```



```
plt.xticks(rotation = 0)
plt.show()
```



```
[308]: #Visualising the target variable
sns.countplot(x = "Survived", data = titanic)
plt.title("Distribution of target variable")
plt.xlabel("Survived (0 = No, 1 = Yes)")
plt.xticks(rotation = 0)
plt.show()
```



```
[309]: #Visualising categorical values
titanic.select_dtypes('object')
```

```
[309]:
```


	Name	Sex	Ticket	Number \
0	Allen, Miss. Elisabeth Walton	female	24160	
1	Allison, Master. Hudson Trevor	male	113781	
2	Allison, Miss. Helen Loraine	female	113781	
3	Allison, Mr. Hudson Joshua Creighton	male	113781	
4	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	female	113781	
...	
1304	Zabour, Miss. Hileni	female	2665	
1305	Zabour, Miss. Thamine	female	2665	
1306	Zakarian, Mr. Mapriededer	male	2656	
1307	Zakarian, Mr. Ortin	male	2670	
1308	Zimmerman, Mr. Leo	male	315082	

	Embarked
0	Southampton
1	Southampton
2	Southampton

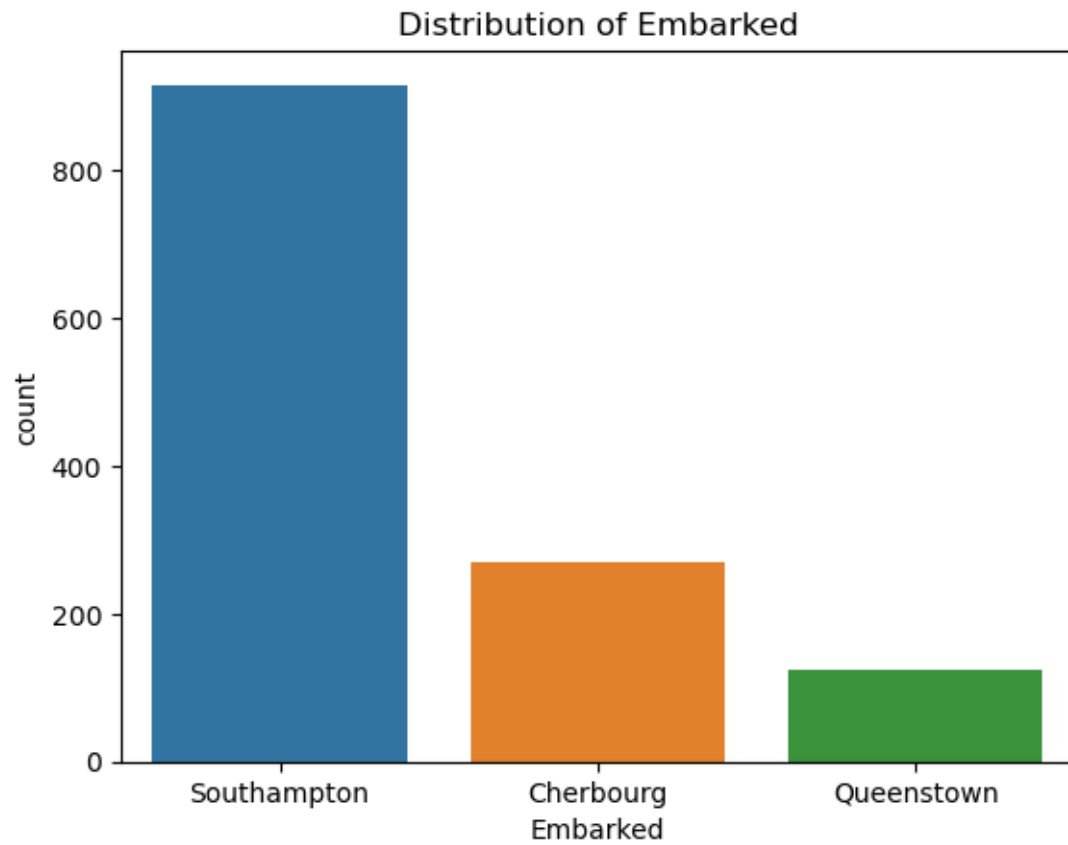
```
3      Southampton
4      Southampton
...
1304    Cherbourg
1305    Cherbourg
1306    Cherbourg
1307    Cherbourg
1308    Southampton
```

```
[1309 rows x 4 columns]
```

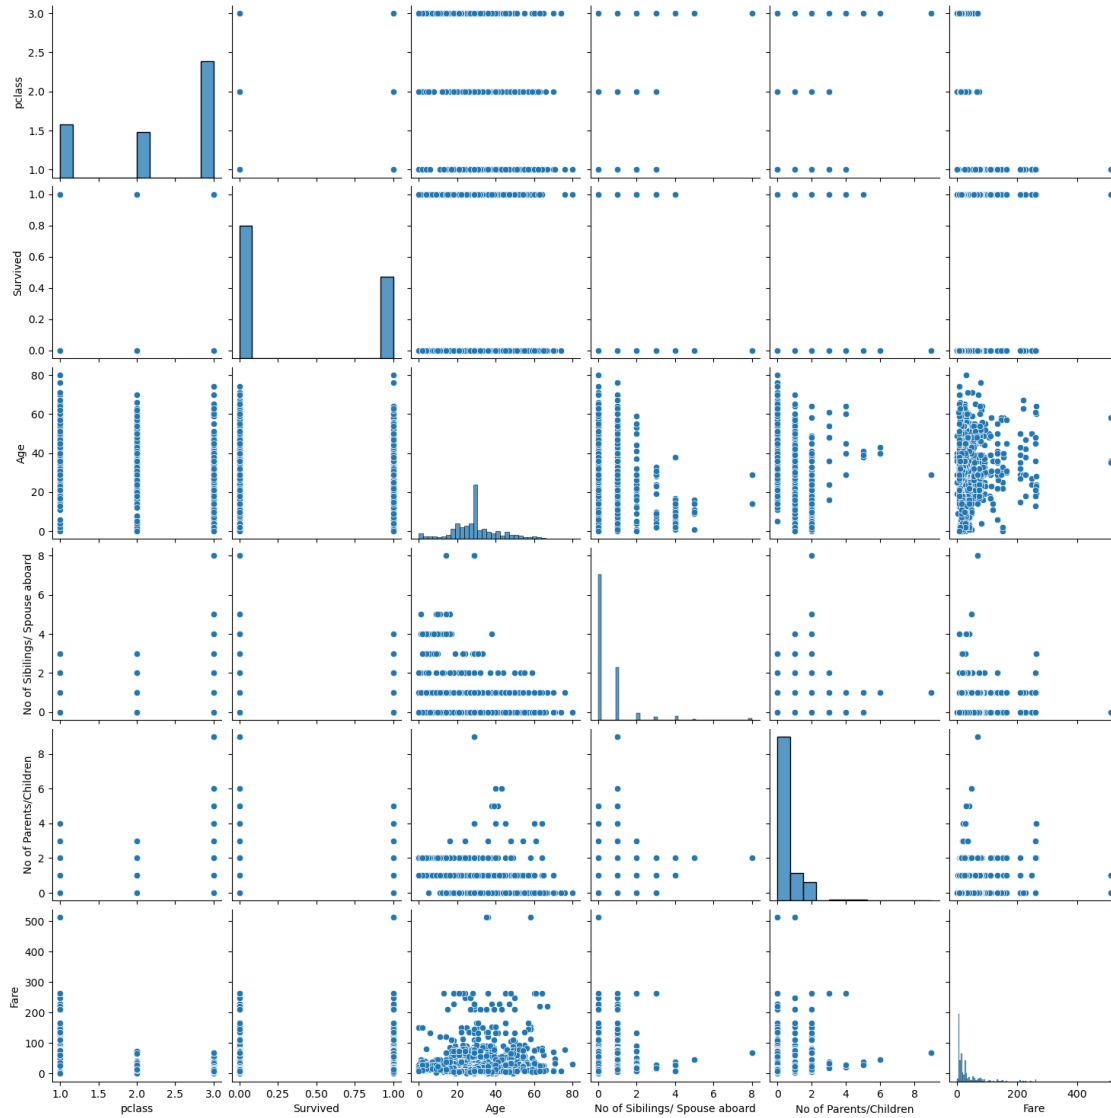
```
[310]: #Visualising Categorical Variables
sns.countplot(x = 'Sex', data = titanic)
plt.title(" Sex Distribution")
plt.show()
```



```
[311]: sns.countplot(x = "Embarked", data = titanic)
plt.title("Distribution of Embarked")
plt.show()
```



```
[312]: sns.pairplot(data = titanic)
plt.show()
```



```
[313]: #Exploring Numerical Values
titanic.select_dtypes('int')
```

```
[313]:
```

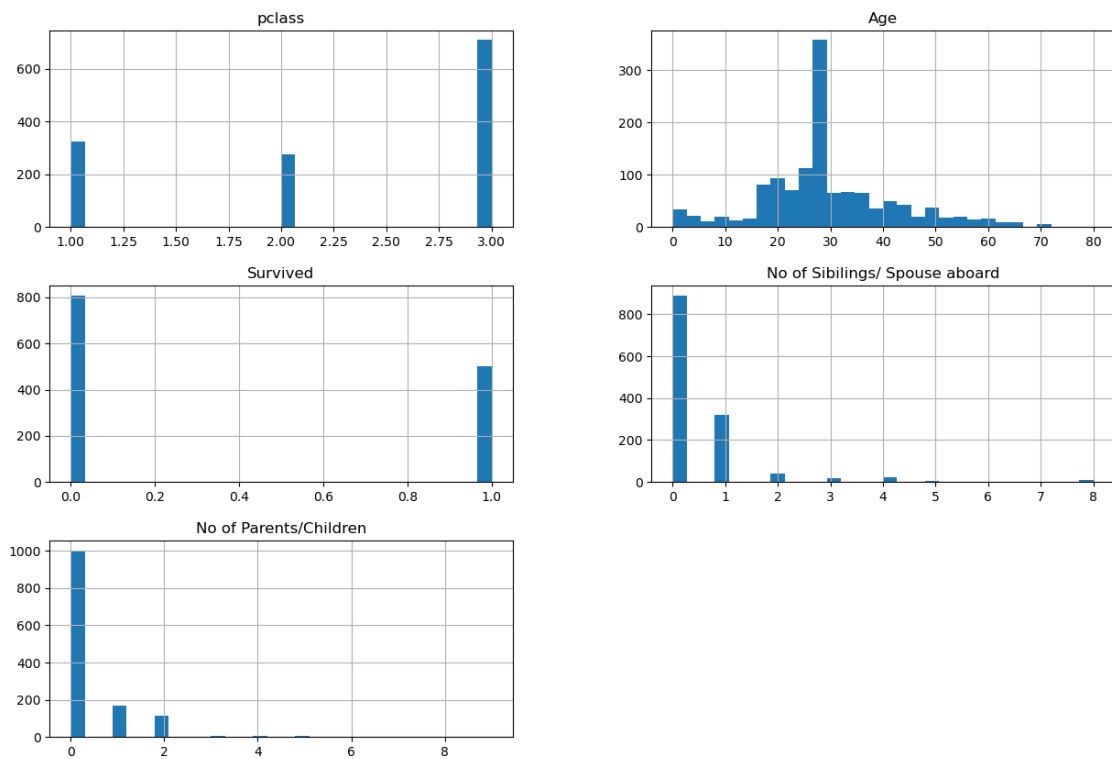
	pclass	Survived	Age	No of Sibilings/ Spouse aboard	\
0	1	1	29		0
1	1	1	0		1
2	1	0	2		1
3	1	0	30		1
4	1	0	25		1
...
1304	3	0	14		1
1305	3	0	29		1
1306	3	0	26		0

1307	3	0	27	0
1308	3	0	29	0

No of Parents/Children		
0		0
1		2
2		2
3		2
4		2
...	...	
1304		0
1305		0
1306		0
1307		0
1308		0

[1309 rows x 5 columns]

```
[314]: #Visualising numerical features
numerical_features = ["pclass","Age","Survived","No of Sibilings/ Spouse_
↪aboard", "No of Parents/Children"]
titanic[numerical_features].hist(bins = 30, figsize= (15,10))
plt.show()
```



Model Preprocessing

```
[315]: #Encoding categorical values
cat_values = ["Name", "Sex", "Ticket Number", "Embarked"]
titanic= pd.get_dummies(titanic, columns = cat_values, drop_first=True)
```

```
[316]: #Label encoding the target variable
le = LabelEncoder()
titanic['Survived']= le.fit_transform( titanic['Survived'])
```

```
[317]: #Feature Scaling
scaler = StandardScaler()
titanic[numerical_features] = scaler.fit_transform(titanic[numerical_features])
```

Model Building

```
[318]: y = titanic["Survived"].info()
titanic["Survived"]=titanic["Survived"].astype("int64")
#y.astype(int)
#y_train.astype(int)
#y_train.info()
```

```
<class 'pandas.core.series.Series'>
Int64Index: 1309 entries, 0 to 1308
Series name: Survived
Non-Null Count  Dtype
-----  -----
1309 non-null   float64
dtypes: float64(1)
memory usage: 20.5 KB
```

```
[319]: #train_test split
X = titanic.drop("Survived", axis = 1)
y = titanic["Survived"]
X_standardised = scaler.fit_transform(X)
X_train,X_test,y_train,y_test = train_test_split(X_standardised,y, test_size= 0.
↪2, random_state=42)
```

```
[320]: #Model Training using LogisticRegression
Lg = LogisticRegression()
Lg.fit(X_train,y_train)
```

```
[320]: LogisticRegression()
```

```
[321]: #Model Training using RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(X_train,y_train)
```

[321]: RandomForestClassifier()

HyperParameter Tuning

```
[322]: #Hyperparameter tuning using gridsearchcv for RandomForest
# Define the hyperparameter grid
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth': [10, 20, 30, None],
    'criterion': ['gini']
}
grid_search = GridSearchCV(estimator= rfc, param_grid= param_grid,cv = 3,
    ↪n_jobs= -1, verbose= 2)
grid_search.fit(X_train, y_train)

# Best hyperparameters
print("Best Hyperparameters:", grid_search.best_params_)

# Best estimator
best_rfc = grid_search.best_estimator_
```

Fitting 3 folds for each of 36 candidates, totalling 108 fits

```
c:\Users\elmaf\anaconda3\Lib\site-
packages\joblib\externals\loky\process_executor.py:700: UserWarning: A worker
stopped while some jobs were given to the executor. This can be caused by a too
short worker timeout or by a memory leak.
```

```
warnings.warn(
c:\Users\elmaf\anaconda3\Lib\site-
packages\sklearn\model_selection\_validation.py:425: FitFailedWarning:
36 fits failed out of a total of 108.
The score on these train-test partitions for these parameters will be set to
nan.
If these failures are not expected, you can try to debug them by setting
error_score='raise'.
```

Below are more details about the failures:

```
-----
36 fits failed with the following error:
Traceback (most recent call last):
  File "c:\Users\elmaf\anaconda3\Lib\site-
packages\sklearn\model_selection\_validation.py", line 732, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\base.py", line 1144,
in wrapper
    estimator._validate_params()
  File "c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\base.py", line 637,
```



```

in _validate_params
    validate_parameter_constraints(
        File "c:\Users\elmaf\anaconda3\Lib\site-
packages\sklearn\utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features'
parameter of RandomForestClassifier must be an int in the range [1, inf), a
float in the range (0.0, 1.0], a str among {'sqrt', 'log2'} or None. Got 'auto'
instead.

warnings.warn(some_fits_failed_message, FitFailedWarning)
c:\Users\elmaf\anaconda3\Lib\site-
packages\sklearn\model_selection\_search.py:976: UserWarning: One or more of the
test scores are non-finite: [          nan          nan          nan 0.69818529
0.70391595 0.69436485
0.63514804 0.63514804 0.63514804          nan          nan          nan
0.77936963 0.78127985 0.78319007 0.63514804 0.63610315 0.63705826
          nan          nan          nan 0.78796562 0.79465138 0.79369628
0.65616046 0.65902579 0.6599809          nan          nan          nan
0.80897803 0.80993314 0.80993314 0.79178606 0.80324737 0.80993314]
warnings.warn(

Best Hyperparameters: {'criterion': 'gini', 'max_depth': None, 'max_features':
'sqrt', 'n_estimators': 200}

```

```

[345]: #Model Predictions
y_pred = best_rfc.predict(X_test)
y_pred

```

```

[345]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0,
0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0,
0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1,
1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1,
1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1,
0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1,
0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0, 1, 0,
1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0,
1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0,
1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1],
dtype=int64)

```

```

[346]: #Calculating accuracy
Accuracy = accuracy_score(y_pred,y_test)
Accuracy

```

[346]: 0.7862595419847328

```
[325]: #Hyperparameter tuning using gridsearchcv for Logistic Regression
# Define the hyperparameter grid
param_grid = {'C': [0.1, 1, 10],
              'solver': ['liblinear', 'saga']}
grid = GridSearchCV(estimator = Lg, param_grid=param_grid, cv=3)
grid.fit(X_train, y_train)
best_model = grid.best_estimator_
```

```
c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\linear_model\_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
  warnings.warn(
c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\linear_model\_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
  warnings.warn(
c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\linear_model\_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
  warnings.warn(
c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\linear_model\_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
  warnings.warn(
c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\linear_model\_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
  warnings.warn(
c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\linear_model\_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
  warnings.warn(
c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\linear_model\_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
  warnings.warn(
c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\linear_model\_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
  warnings.warn(
c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\linear_model\_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
  warnings.warn(
c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\linear_model\_sag.py:350:
```

```
ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
```

```
[347]: #Model Predictions
y_pred = best_rfc.predict(X_test)
y_pred
```

```
[347]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0,
        0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0,
        0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1,
        1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1,
        1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1,
        0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1,
        0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0, 1, 0,
        1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0,
        1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0,
        1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1],
        dtype=int64)
```

```
[348]: #Calculating accuracy
Accuracy = accuracy_score(y_pred,y_test)
Accuracy
```

```
[348]: 0.7862595419847328
```

```
[342]: #Hyperparameter tuning using Randomizedsearchcv for RandomForestClassifier
param_grid = {
    'n_estimators': [100, 200,300],
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth': [10, 20, 30, None],
    'criterion': ['gini']
}
# Setup random hyperparameter search for RandomForestClassifier
rs_rf = RandomizedSearchCV(RandomForestClassifier(),
    param_distributions=param_grid, cv=5, n_iter=20, verbose=True)
rs_rf.fit(X_train,y_train)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

c:\Users\elmaf\anaconda3\Lib\site-

packages\sklearn\model_selection_validation.py:425: FitFailedWarning:

40 fits failed out of a total of 100.

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting

```
error_score='raise'.
```

Below are more details about the failures:

```
-----
40 fits failed with the following error:
Traceback (most recent call last):
  File "c:\Users\elmaf\anaconda3\Lib\site-
packages\sklearn\model_selection\_validation.py", line 732, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\base.py", line 1144,
in wrapper
    estimator._validate_params()
  File "c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\base.py", line 637,
in _validate_params
    validate_parameter_constraints(
  File "c:\Users\elmaf\anaconda3\Lib\site-
packages\sklearn\utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features'
parameter of RandomForestClassifier must be an int in the range [1, inf), a
float in the range (0.0, 1.0], a str among {'sqrt', 'log2'} or None. Got 'auto'
instead.
```

```
    warnings.warn(some_fits_failed_message, FitFailedWarning)
c:\Users\elmaf\anaconda3\Lib\site-
packages\sklearn\model_selection\_search.py:976: UserWarning: One or more of the
test scores are non-finite: [0.63515152 0.63610845          nan          nan
0.63515152 0.63515152
0.81564821          nan          nan          nan 0.81850535 0.77840966
          nan 0.7765049 0.6361039          nan          nan 0.65138756
0.81181135 0.80226475]
    warnings.warn(
```

```
[342]: RandomizedSearchCV(cv=5, estimator=RandomForestClassifier(), n_iter=20,
                        param_distributions={'criterion': ['gini'],
                                           'max_depth': [10, 20, 30, None],
                                           'max_features': ['auto', 'sqrt',
                                                           'log2'],
                                           'n_estimators': [100, 200, 300]},
                        verbose=True)
```

```
[349]: #Model Evaluation using RandomisedSearchCV - RandomForestClassifier
y_pred_rs_rf = rs_rf .predict(X_test)
y_pred_rs_rf
```

```
[349]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0,
          0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0,
          0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1,
          1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1,
          1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1,
          0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1,
          0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0, 1, 0,
          1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0,
          1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1],
          dtype=int64)
```

```
[350]: #Model Accuracy
Accuracy_rnd = accuracy_score(y_pred_rs_rf,y_test)
Accuracy_rnd
```

```
[350]: 0.7900763358778626
```

```
[329]: #Hyperparameter tuning using Randomizedsearchcv for Logistic Regression
# Define the hyperparameter grid
param_grid = {'C': [0.1, 1, 10],
              'solver': ['liblinear', 'saga']}
# Initialize RandomizedSearchCV
rs_log_reg = RandomizedSearchCV(estimator=LogisticRegression(),
    ↪param_distributions=param_grid,
                                cv=3, n_iter=20, verbose=True, random_state=42,
    ↪n_jobs=-1)

# Fit RandomizedSearchCV to the data
rs_log_reg.fit(X_train, y_train)
```

```
c:\Users\elmaf\anaconda3\Lib\site-
packages\sklearn\model_selection\_search.py:307: UserWarning: The total space of
parameters 6 is smaller than n_iter=20. Running 6 iterations. For exhaustive
searches, use GridSearchCV.
```

```
warnings.warn(
```

```
Fitting 3 folds for each of 6 candidates, totalling 18 fits
```

```
c:\Users\elmaf\anaconda3\Lib\site-packages\sklearn\linear_model\_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
```

```
warnings.warn(
```

```
[329]: RandomizedSearchCV(cv=3, estimator=LogisticRegression(), n_iter=20, n_jobs=-1,
          param_distributions={'C': [0.1, 1, 10],
```

```
'solver': ['liblinear', 'saga']},
random_state=42, verbose=True)
```

Model Evaluation

```
[351]: #Model Evaluation using RandomisedSearchCV - Logistic Regression
y_pred_rcv = rs_log_reg.predict(X_test)
y_pred_rcv
```

```
[351]: array([0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0,
0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0,
0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1,
1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1,
1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1,
0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1,
0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0,
0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1,
1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0,
1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1],
dtype=int64)
```

```
[352]: #Model Accuracy
Accuracy_rnd_reg= accuracy_score(y_pred_rcv ,y_test)
Accuracy_rnd_reg
```

```
[352]: 0.7748091603053435
```

From the 2 models used Logistic Regression and RandomForest Classifier, the best performing model is RandomForestClassifier model hypertuned using RandomisedSearchCV


The model is 79% accurate

```
[361]: # Evaluation metrics
print(confusion_matrix(y_test, y_pred_rs_rf))
print(classification_report(y_test, y_pred_rs_rf))
sns.heatmap(confusion_matrix(y_test, y_pred_rs_rf), annot=True, fmt='d', cmap=
↳ 'Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.xticks([0, 1], ['Predicted Negative', 'Predicted Positive'])
plt.yticks([0, 1], ['Actual Negative', 'Actual Positive'])
plt.show()
```

```
[[137   7]
 [ 48  70]]

precision    recall  f1-score   support
```

	0	0.74	0.95	0.83	144
	1	0.91	0.59	0.72	118
accuracy				0.79	262
macro avg		0.82	0.77	0.78	262
weighted avg		0.82	0.79	0.78	262




Explanation: True Positives (TP): Correctly predicted positive cases.

True Negatives (TN): Correctly predicted negative cases.

False Positives (FP): Incorrectly predicted positive cases. False Negatives (FN): Incorrectly predicted negative cases.

```
[359]: #ROC display curve
RocCurveDisplay.from_estimator(estimator= rs_rf,X = X_test, y =y_test);
plt.title("Receiver Operating Characteristic (ROC) curve")
```

```
[359]: Text(0.5, 1.0, 'Receiver Operating Characteristic (ROC) curve')
```

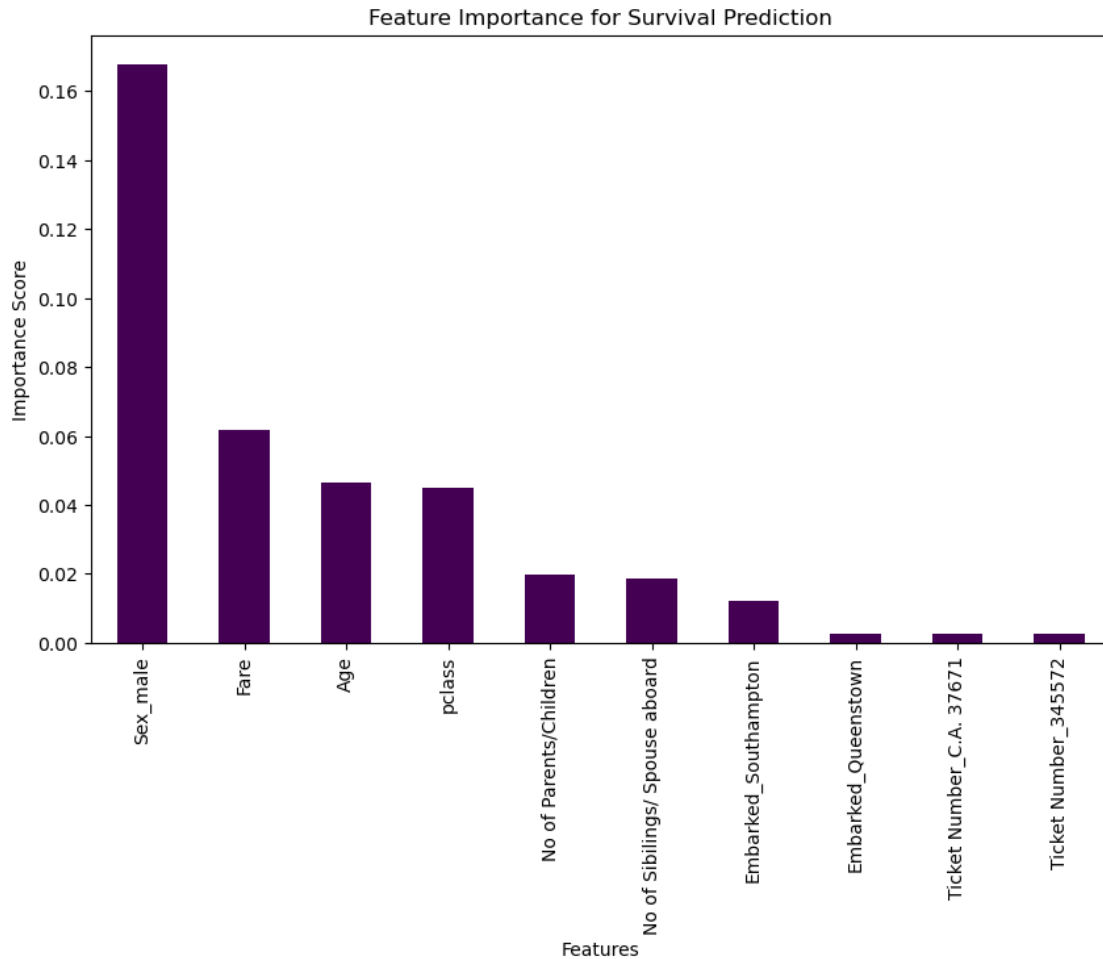


This is great, our model does far better with an AUC of 0.86, but a perfect model would achieve an AUC score of 1.0, so there's still room for improvement.

```
[379]: #Feature Importance
best_rf = rs_rf.best_estimator_ # Retrieve the best estimator from
↳ RandomizedSearchCV
features = X.columns

# Get feature importances from the best RandomForestClassifier model
feature_importances = best_rf.feature_importances_
feature_importances = pd.Series(feature_importances, index=features)
feature_importances = feature_importances.sort_values(ascending=False).head(10)
plt.figure(figsize=(10, 6))
feature_importances.plot(kind='bar', cmap='viridis')
plt.title('Feature Importance for Survival Prediction')
plt.xlabel('Features')
plt.ylabel('Importance Score')
```

```
[379]: Text(0, 0.5, 'Importance Score')
```

The bar plot shows the importance of each feature in making predictions. Higher importance indicates a greater influence on the model's predictions.

Conclusion

Our model is 79% accurate. It is fairly reasonable but there is room for improvement

Key Findings

Gender is the most significant predictor of survival followed by class and age

Female passengers had a higher chance of surviving compared to male passengers